

Managing and Controlling the Kubernetes Scheduler



Kien Bui

DevOps & Platform Engineer

Course Overview



Configuring and Managing Storage in
Kubernetes

Configuration as Data -Environment Variables,
Secrets and ConfigMaps

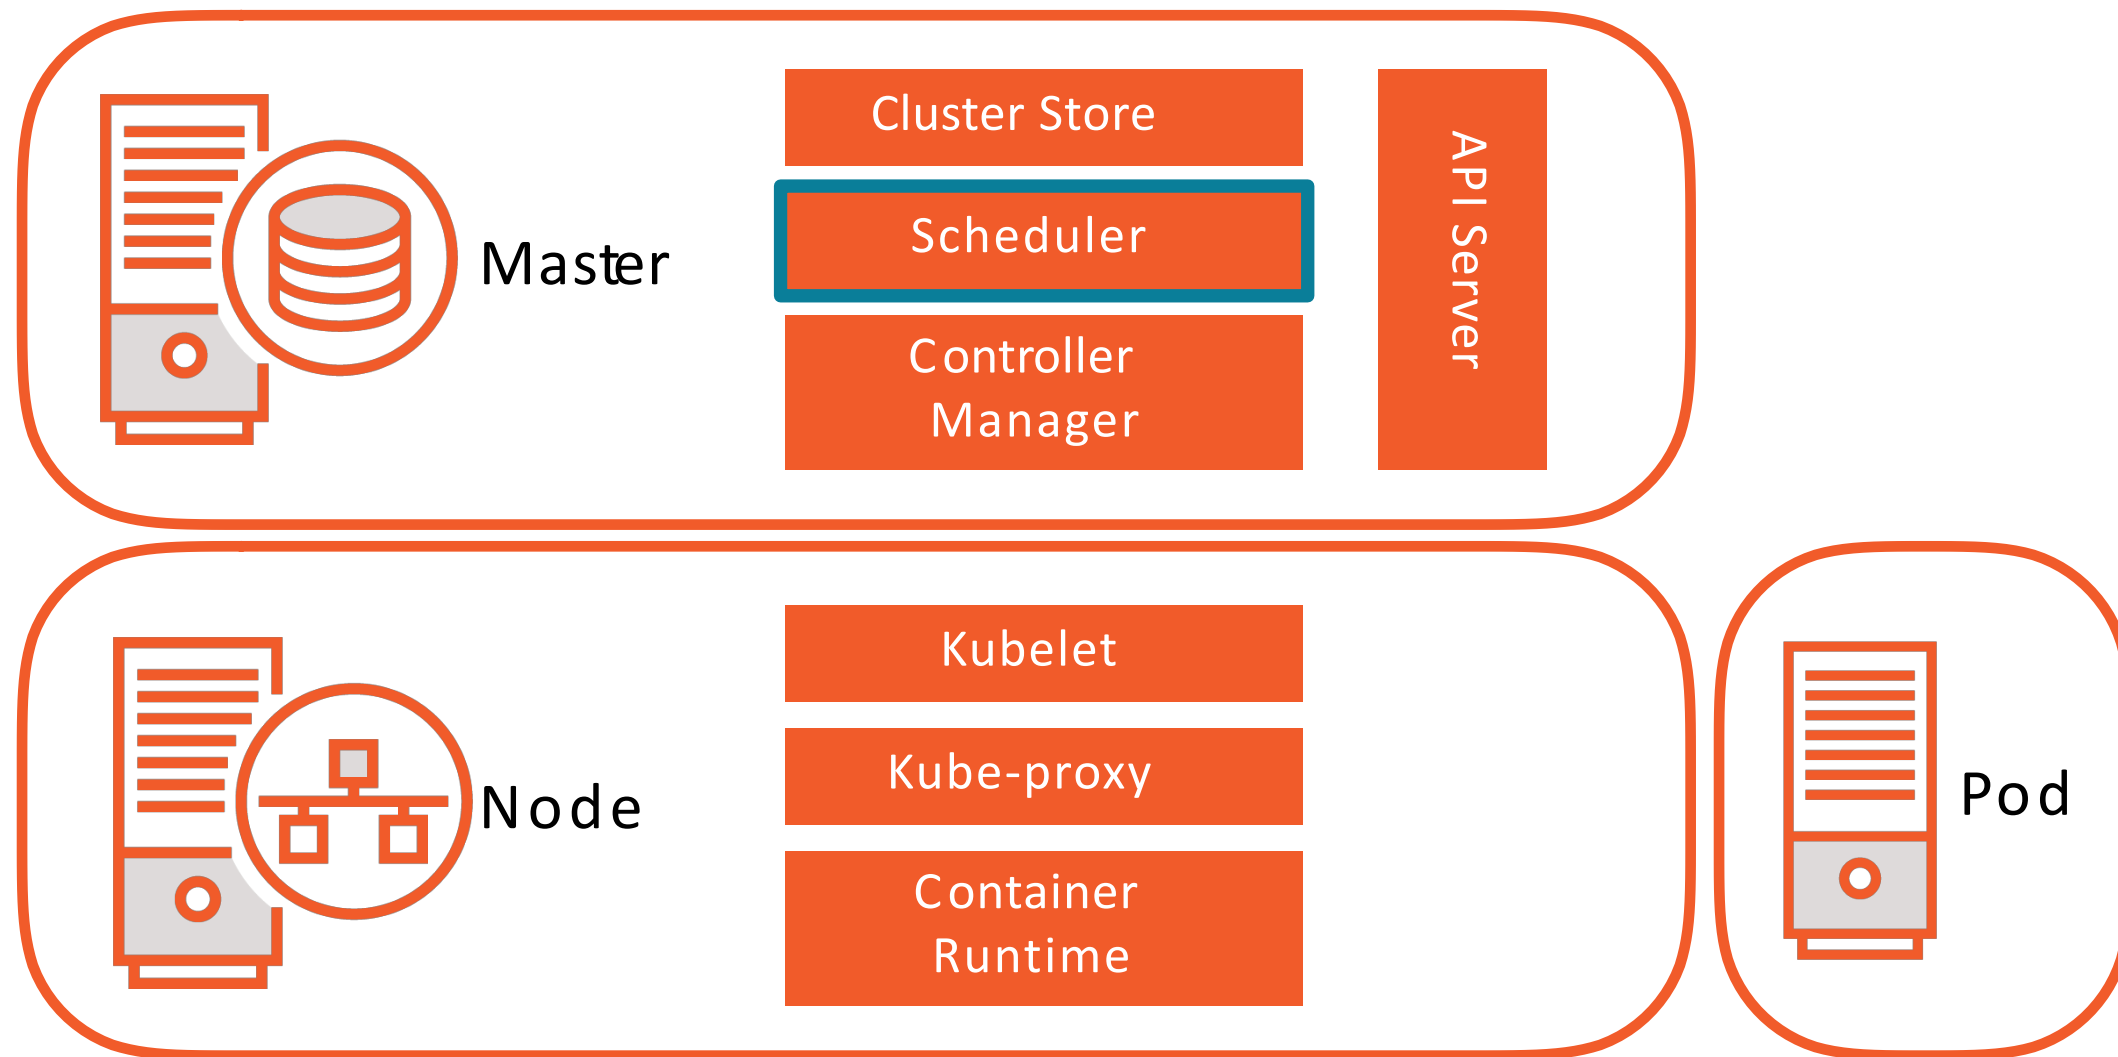
Managing and Controlling the
Kubernetes Scheduler

Overview

Scheduling in Kubernetes


Controlling scheduling in Kubernetes

Master Node - Control Plane



Kubernetes has one job...
starting Pods on Nodes

Scheduling in Kubernetes



Selecting a Node to start a Pod on
`kube-scheduler`

Scheduling in Kubernetes



Resources



Policy

Scheduling Process

Watches the API
Server for
Unscheduled Pods

Node selection

Update nodeName in
the Pod object

Nodes' kubelets
watch API Server
for work

Signal container
runtime to start
container(s)

Node Selection

Filtering

From all Nodes

Apply Filters

Filtered Nodes

Hard constraints

Scoring

Scoring functions

Feasible Nodes

Policy constraints

Binding

Selected Nodes List

Ties are broken


Update API Object

nodeName: c1-node1



Pod

Resource Requests



Setting `requests` will cause the scheduler to find a Node to fit the workload/Pod

`requests` are guarantees

CPU

Memory

Allocatable resources per Node

Pods that need to be scheduled but there not enough resources available will go `Pending`

Demo

Scheduling in action

Scheduling Pods with `requests`

Controlling Scheduling

Node Selector

Affinity

Taint and
Tolerations

Node Cordoning

Manual Scheduling

Node Selector



`nodeSelector` -assign Pods to Nodes using Labels and Selectors

Apply Labels to Nodes

Scheduler will assign Pods to a Node with a matching Label

Simple key/value check based on `matchLabels`


Often used to map Pods to Nodes based on...

- Special hardware requirements

- Workload isolation

Managing Kubernetes API Server and Pods

Scheduling - Node Selector



Assigning Pods to Nodes using Node Selectors

```
kubectl label node c1-node3 hardware=local_gpu
```

```
spec:
```

```
  containers:
```

```
  - name: hello-world
```

```
    image: gcr.io/google-samples/hello-app:1.0
```

```
    ports:
```

```
    - containerPort: 8080
```

```
  nodeSelector:
```

```
    hardware: local_gpu
```

Affinity and Anti-Affinity



`nodeAffinity` - uses Labels on Nodes to make a scheduling decision with `matchExpressions`

`requiredDuringSchedulingIgnoredDuringExecution`

`preferredDuringSchedulingIgnoredDuringExecution`


`podAffinity` - schedule Pods onto the same Node, Zone as some other Pod

`podAntiAffinity` - schedule Pods onto the different Node, Zone as some other Pod

Managing Kubernetes API Server and Pods

<https://kubernetes.io/docs/concepts/configuration/assign-pod-node/#affinity-and-anti-affinity>

Scheduling - Pod Affinity/Anti-Affinity



Using Affinity to Control Pod Placement

```
spec:
  containers:
  - name: hello-world-cache
  ...
  affinity:
    podAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
      - labelSelector:
          matchExpressions:
          - key: app
            operator: In
            values:
            - hello-world-web
        topologyKey: "kubernetes.io/hostname"
```

Taints and Tolerations



Taints - ability to control which Pods are scheduled to Nodes

Tolerations - allows a Pod to ignore a Taint and be scheduled as normal on Tainted Nodes

Useful in scenarios where the cluster administrator needs to influence scheduling without depending on the user

```
key=value:effect
```

```
kubectl taint nodes c1-node1 \
  key=MyTaint:NoSchedule
```

Scheduling - Taints and Tolerations



Adding a Taint to a Nodes and a Toleration to a Pod

```
kubectl taint nodes c1-node1 key=MyTaint:NoSchedule
```

```
spec:
```

```
  containers:
```

```
  - name: hello-world
```

```
    image: gcr.io/google-samples/hello-app:1.0
```

```
    ports:
```

```
    - containerPort: 8080
```

```
  tolerations:
```

```
  - key: "key"
```

```
    operator: "Equal"
```

```
    value: "MyTaint"
```

```
    effect: "NoSchedule"
```

Demo

Using Affinity and Anti-Affinity to schedule
Pods to Nodes

Controlling Pod placement with Taints and
Tolerations

Node Cordoning



Marks a Node as unschedulable

Prevents new Pods from being scheduled to that Node

Does not affect any existing Pods on the Node

This is useful as a preparatory step before a Node reboot or maintenance

```
kubectl cordon c1-node3
```

If you want to gracefully evict your Pods from a Node...

```
kubectl drain c1-node3 --ignore-daemonsets
```

Manually Scheduling a Pod



Scheduler populates `nodeName`

If you specify `nodeName` in your Pod definition the Pod will be started on that node

Node's name must exist

Still subject to Node resource constraints

Configuring Multiple Schedulers



Implement your own scheduler

Run multiple schedulers concurrently

Define in your Pod Spec which scheduler you want

Deploy your scheduler as a system Pod in the cluster

<https://kubernetes.io/docs/tasks/administer-cluster/configure-multiple-schedulers/>

Demo

Node Cordoning

Manually scheduling a Pod

Review

Scheduling in Kubernetes

Controlling scheduling in Kubernetes